

Recommended Multi-Agent Chat Architecture for Autogen v0.4 Chatbot

Project Overview and Requirements

You are building a command-line **multi-agent chatbot** using Microsoft Autogen v0.4. The chatbot must guide a user through a **login/registration workflow** with the following steps:

1. **Username Input & Validation:** Prompt the user for a username and validate the input format.
2. **User Lookup:** Check if the username exists in a JSON database.
3. **Handle Missing/Invalid Users:** If the username is invalid or not found, handle errors or initiate a new user registration flow.
4. **New User Registration:** If needed, collect and parse additional user details and create a new user entry (with proper validation).
5. **LinkedIn Login Verification:** Use an automated browser (Playwright) to verify the user's LinkedIn credentials.
6. **Credential Storage:** Encrypt and store user credentials after successful login.
7. **Session Management:** Save login session information for future use.

Design Criteria: The system should follow a mostly **predefined sequence** of steps (as above) but allow some **dynamic branching** based on context (e.g. existing vs. new user). Multiple specialized agents need to collaborate with a degree of independence while sharing context. The solution should be **simple yet scalable** – easy to implement incrementally and debug, but also extensible to future features (such as resume parsing or employer profiling modules). Human oversight is important: the framework should allow a **human-in-the-loop** (e.g. to confirm actions or provide missing info) when necessary. Long-term memory or state persistence between runs is a nice-to-have (for remembering user data or conversation history), but not a strict requirement.

Given these requirements, we will compare Autogen v0.4's multi-agent architectures and recommend the best approach.

Autogen v0.4 Multi-Agent Architecture Options

Microsoft Autogen v0.4 provides several patterns for organizing teams of agents. The main choices you've identified are: **RoundRobinGroupChat**, **SelectorGroupChat**, **Magentic-One**, **GraphFlow**, and building a **custom solution**. Below is an overview of each, with an analysis of how they align with your needs:

RoundRobinGroupChat

Overview: RoundRobinGroupChat is the simplest team configuration. Agents take turns speaking in a fixed round-robin order, **sharing a single conversation context** ¹. Each agent's message is broadcast to all others, ensuring everyone sees the same context. In effect, the agents form a "panel" that responds in sequence: agent 1, then agent 2, then agent 3, and so on (wrapping around continuously).

Pros: This pattern is **easy to implement and reason about**. All agents operate on the same conversation state, and the fixed turn order guarantees predictability ¹. RoundRobin is excellent for straightforward scenarios where agents should simply alternate (for example, an AI assistant and a user proxy trading turns). It provides a **consistent shared context** without complex coordination logic ¹.

Cons: RoundRobinGroupChat lacks flexibility for dynamic flows. The strict alternation makes it difficult to **skip or re-order agents** based on context. Every agent gets called in turn regardless of whether its action is needed at that step. Conditional branching (e.g. only invoke a “registration agent” if the user is new) is clumsy to implement – you would have to either include logic inside the agent’s prompt to sometimes do nothing, or end the round-robin loop early, which can be awkward. In short, RoundRobin handles **linear, predetermined sequences** well, but it does not naturally support making decisions about *which* agent should act next. This could become a limitation in your workflow, where certain steps (like registration) should only occur in specific conditions.

Use Case Fit: If your flow were strictly linear for all users, RoundRobin could work (e.g. always: Ask user → Validate → Check DB → Login → Done, cycling through a fixed set of agent roles). However, because you have **conditional branches** (existing vs new user, valid vs invalid input) and possibly loops (retrying input until valid), a pure round-robin sequence would become cumbersome. It’s still possible to prototype with RoundRobin for very simple interactions, but it may not scale cleanly to the full logic you need.

SelectorGroupChat

Overview: SelectorGroupChat is a more **dynamic team** pattern. Like RoundRobin, all agents share the conversation, but instead of a fixed order, a **generative model selects the next speaker** based on the dialogue context ². Essentially, after each message, Autogen uses an LLM (or a custom function) to decide which agent should speak next. Agents still broadcast their outputs to all others, maintaining a shared context, but the turn-taking is context-sensitive rather than predetermined.

Pros: The SelectorGroupChat enables **context-aware collaboration** ². By default, an LLM “moderator” looks at the conversation so far (and agents’ roles/descriptions) and chooses the most appropriate next agent to respond ³. This allows more flexibility in **branching**: for example, if the conversation indicates the user is new, the model could choose a “RegistrationAgent” to speak next; if not, it might skip to a “LoginAgent”. It also prevents irrelevant agents from chiming in when not needed. Autogen’s SelectorGroupChat comes with features to fine-tune this behavior, such as preventing one agent from dominating by taking consecutive turns (unless allowed) ⁴. Importantly, you can override the selection logic with a **custom selector function** for full control ⁵ – meaning you can programmatically decide the next agent based on the conversation state, instead of relying on the LLM’s choice.

Cons: The added flexibility comes at the cost of complexity. If you use the default LLM-based selection, you’ll need to carefully engineer the system prompt or agent descriptions so the model reliably picks the correct agent at each juncture. This can be unpredictable (especially as new agents or behaviors are added). On the other hand, if you use a custom selection function to deterministically route turns, you are essentially writing your own controller logic – at that point, you might wonder if a more structured approach (like GraphFlow) is cleaner. **Debugging** a model-driven selector can also be tricky, because mis-routings might be due to subtle prompt issues. SelectorGroupChat still operates as a free-flowing conversation; it doesn’t inherently enforce an overall **workflow structure** beyond what the model or your custom function infer. For example, handling a loop (ask again if input invalid) or complex branching logic in a selector function could become as involved as coding a state machine.

Use Case Fit: SelectorGroupChat can handle **semi-structured dialogs** well. If you anticipate the conversation might meander or involve open-ended sub-dialogues, the model-based selection gives a flexible, **adaptive** coordination. In your case, however, the interaction flow is mostly known upfront (it's more a guided procedure than an open-ended collaboration). You likely don't need the full generality of an LLM deciding turn-taking on the fly. It could be leveraged for minor variations (e.g. skipping an "ask LinkedIn credentials" agent if already logged in), but those can also be handled with explicit logic. In short, SelectorGroupChat is powerful for dynamic multi-agent chats, but for a **well-defined sequence with clear decision points**, you might prefer an approach that explicitly encodes that workflow (for transparency and easier debugging).

Magentic-One (Orchestrator Pattern)

Overview: Magentic-One is a ready-made, generalist multi-agent system introduced with Autogen. It is essentially an **orchestrator-based** architecture: a central "Orchestrator" agent dynamically plans tasks and delegates them to a team of specialist agents (web surfer, file surfer, coder, etc.), handling open-ended goals ⁶. This system was designed to solve broad, **open-ended tasks across various domains** (research, web browsing, coding, etc.) autonomously, as described in the technical report ⁶. In Autogen v0.4, Magentic-One is implemented on top of the AgentChat API – it provides a class `MagenticOneGroupChat` (the orchestrator team) and several premade agent types (e.g. `MultimodalWebSurfer`, `FileSurfer`, etc.) ⁷.

Example of a complex multi-agent orchestration (Magentic-One): The central Orchestrator agent creates a plan, delegates sub-tasks to specialized agents (for web browsing, file reading, coding, etc.), and adjusts the plan as tasks are completed ⁸. This showcases how an agent can coordinate a team to achieve a high-level goal.

Pros: Magentic-One represents a **state-of-the-art orchestration approach**. The orchestrator agent (backed by a powerful LLM) can break down a complex objective into steps, assign those steps to the appropriate agents, and integrate the results, all while adjusting dynamically to new information ⁸. It achieves strong performance on agent benchmarks by **dynamically revising plans** and handling unexpected outcomes autonomously ⁸. For a developer, Magentic-One offers a sort of blueprint of how to structure a **general problem-solving agent team** – you could potentially adapt its orchestrator concept to your own needs. It's also modular: since it's built on AgentChat, you can incorporate any custom AgentChat agents or tools into the orchestrator's arsenal ⁷.

Cons: The power of Magentic-One comes with **significant complexity**. It is essentially an AI **planner** in the loop, which might be overkill for a relatively routine, predetermined workflow. Tuning an autonomous planner agent (ensuring it understands your specific task sequence and doesn't go off track) could be as challenging as just coding the sequence directly. Moreover, using Magentic-One "as-is" means yielding a lot of control to the LLM orchestrator – fine for open tasks, but perhaps unnecessary if you already know the exact steps. It's also a newer solution; while it's designed to be user-friendly in v0.4, it may involve more moving parts (multiple agent classes, tools, etc.) than a purpose-built flow. In short, **Magentic-One is a heavy-weight, generalized solution** – likely more complex than needed for your login chatbot. Adopting it would mean framing your problem as a planning task for the orchestrator agent, which might complicate incremental development and debugging.

Use Case Fit: For your scenario (user login/registration), Magentic-One is not an ideal fit. The task is *procedural* and domain-specific, not an open-ended problem that requires dynamic goal decomposition. While you can certainly learn from the orchestrator pattern, a simpler architecture will give you more direct control. Magentic-One shines in scenarios where you want the AI to autonomously figure out *how*

to accomplish a goal by calling various agents – here, you already know the steps, so a lighter-weight coordination method is preferable. (That said, as your project grows to include more complex modules like resume parsing or job research, you could consider incorporating some planning elements, but it would still likely be overkill to deploy a full Magentic-One orchestrator for each user session.)

GraphFlow (Directed Workflow)

Overview: GraphFlow is a **structured workflow engine** for Autogen's AgentChat. It allows you to explicitly define a directed graph (nodes and edges) that dictates how agents will interact and in what order ⁹ ¹⁰. Each node in the graph is essentially an agent execution step, and edges represent possible transitions to the next agent. You can include **conditions on edges**, enabling branching logic based on content of agent messages or other state, as well as support loops and even parallelism ¹¹ ¹². In simpler terms, GraphFlow lets you design a **state machine or flow chart** for your multi-agent conversation.

Pros: GraphFlow offers **precise control** over the conversation sequence. You literally draw out the flow of which agent acts after which, including conditional branches for different outcomes. This makes it ideal when you have a **deterministic or semi-deterministic process** (like a login workflow) – you can be sure the agents execute in the intended order and handle each decision point explicitly. The framework supports sequential chains, conditional branching (if/else logic), loops with exit conditions (for retries or iterative tasks), and even parallel steps if needed ¹². The Autogen docs recommend using GraphFlow **when you need strict control over agent ordering or when different outcomes lead to different next steps** ¹³. In other words, if ad-hoc conversation is not enough and you require *deterministic, complex multi-step processes with possible cycles*, GraphFlow is the go-to ¹³. Another benefit is **observability**: since the flow is defined, it's easier to trace which path was taken. Autogen provides tooling to observe and debug flows (e.g. you can inspect node-by-node execution and use the `Console` to stream events), which is great for step-by-step debugging of your agent workflow.

Cons: GraphFlow is a relatively **advanced feature (experimental in v0.4)** ¹⁴. Setting up the directed graph requires more upfront coding (using the `DiGraphBuilder` or similar) compared to using a simple RoundRobin or Selector. It's essentially a form of programming the conversation logic, so it demands clarity on all the possible branches and outcomes. If your process changes, you need to update the graph accordingly (less flexibility for truly unstructured conversation). Also, being new, you might encounter a learning curve or minor API changes as it evolves. However, since you desire control and have a clear idea of the flow, this "con" is more about effort than capability. Once implemented, the graph should actually simplify handling complex paths that would be hard to manage with RoundRobin/Selector alone.

Use Case Fit: GraphFlow appears to be **highly suited to your scenario**. The user onboarding/login process can be naturally mapped to a flowchart (a directed graph of steps). You have distinct phases (input, validation, branch on user existence, etc.) that fit well into a graph model with conditional edges. For example, you can create a node for "Validate Username Agent" -> then an edge to either "Lookup Agent" (if input was valid) or back to "Validate" (if invalid, to loop asking again). The "Lookup Agent" node could have an edge to "Registration Agent" if user not found, or to "Login Agent" if user exists. This is straightforward to encode with GraphFlow's conditional transitions. By using GraphFlow, you'll get a **clear, enforceable sequence** – agents won't talk out of turn or unexpectedly, they will only act when they are activated in the graph. It directly addresses your need for a **mostly predefined flow with some dynamic branches**. While GraphFlow may require more initial setup than the simpler GroupChat patterns, it offers confidence that as you scale to future modules (additional nodes/branches), the complexity remains manageable. The Autogen team specifically notes that one should *"start with a simple team (RoundRobin/Selector) if ad-hoc flow is sufficient, and transition to GraphFlow when*

you need deterministic control or complex branching”¹³ – in your case, we already know you need branching and determinism, so it’s likely time to use GraphFlow.

Custom Agents or Custom Coordination Logic

Overview: Autogen v0.4 is designed to be modular and extensible. If none of the provided patterns fully satisfy your needs, you can always **customize at a lower level**. This could mean writing a custom team coordinator by subclassing the base `GroupChat` or using the core message-passing API, or implementing custom agent classes beyond the built-ins. The Autogen documentation emphasizes that you can “**customize to your need with the base GroupChat in the core**” if RoundRobin/Selector presets don’t fit¹⁵. Similarly, creating a **Custom Agent** class is recommended when the preset agents (AssistantAgent, UserProxyAgent, etc.) don’t do what you need¹⁶.

Pros: Building a custom solution gives you **maximum flexibility**. You are not constrained by the patterns of RoundRobin, Selector, etc., and can implement any interaction protocol or agent behavior you want. For example, you could create a bespoke orchestrator agent that precisely coordinates others (essentially your own lightweight Magentic-One), or an agent that directly calls your JSON database or Playwright script without needing an LLM prompt. If your use case had truly unique requirements, custom coding could tailor the system exactly.

Cons: The obvious downside is **complexity and effort**. You’ll be re-inventing coordination logic that the Autogen frameworks already provide and test. Unless you have a very special case, using a provided architecture (or GraphFlow with conditions) is easier and less error-prone than writing your own multi-agent loop or state machine from scratch. Custom agents are useful for integrating specific tools or APIs (for instance, you might make a “LinkedInAgent” that encapsulates Playwright-based login steps), but these can often be accomplished by using Autogen’s existing agent types with tool integration. In short, while you *can* drop down to the core API or custom classes, you should do so **only if absolutely necessary**. It’s usually better to leverage the framework’s structure (and just plug in custom code at points) than to build an entire coordination mechanism manually.

Use Case Fit: You likely do not need to implement a completely custom team coordinator given the options above. GraphFlow or a Selector with custom function can already handle the branching and sequence you require. However, you *may* implement **custom agents** for certain tasks: for example, a custom code-execution agent that runs Python functions (though Autogen has `CodeExecutorAgent` for this), or a specialized agent that interfaces with the LinkedIn login (possibly by calling a Playwright script internally). These custom agents can be plugged into the GraphFlow or Selector teams. So, keep the door open for custom agent classes (for tool integration or enforcing particular behavior), but you shouldn’t need to replace the overall architecture with a custom one. Only consider writing a custom coordination loop if you hit limitations with the provided frameworks.

Recommended Architecture and Rationale

Considering the above options and your project’s requirements, **the best-suited architecture is to use Autogen’s GraphFlow** to orchestrate your agents through the login workflow, possibly combined with a **UserProxyAgent for human-in-the-loop inputs**. This approach offers a balance of **structured control and modularity** that aligns with your needs. Below is the reasoning:

- **Predefined Flow with Conditional Branching:** GraphFlow directly models your predetermined sequence with branching logic. You can encode the exact workflow (validating input, checking DB, branching to registration or login) as a directed graph. This gives you *deterministic execution*:

each step will happen in order, and each decision (existing vs new user) will lead down a defined path ¹³. This is far clearer than trying to coax a round-robin or selector to follow the same logic implicitly. If the user input is invalid, the graph can loop back to the username prompt node, ensuring the agent will ask again – something that is natural in a graph, but awkward in a free-form chat loop. Essentially, GraphFlow serves as an explicit **workflow diagram** for your agents, which is exactly what a scripted process needs.

- **Agent Collaboration with Independence:** In GraphFlow, each node is an agent performing its role. Agents still have the freedom to “do their job” at their node (e.g., the DB-check agent will decide how to query the JSON and return result), but they won’t interfere with each other out of turn. All agents can share memory/context as needed, but because GraphFlow triggers them one by one (or in controlled parallel), you avoid chaos. This satisfies the need for clear collaboration: you orchestrate which agent should act when, but within its turn the agent can autonomously use its tools or LLM reasoning to complete the task. The **shared context** feature of teams remains – since GraphFlow is built as a Team, all agents can see messages broadcast on the graph execution. For example, the “LinkedInLoginAgent” can be configured to see the prior conversation (such as the user’s name and credentials provided) as context. This means you get the benefits of a multi-agent conversation (shared information) without losing control of the sequence.
- **Dynamic Intent Handling:** While GraphFlow is structured, it can handle dynamic branches through conditional edges. If there are multiple possible outcomes at a step, you can define different next nodes for each. For instance, after attempting LinkedIn login, you might have two outcomes: success (then go to session-saving node) or failure (then maybe go to a retry or error-handling node). You can attach a condition (perhaps based on a message content like “Login success” vs “Login failed”) to choose the appropriate path. This covers “dynamic intent” in the sense of different conversation trajectories based on prior results. Now, if by dynamic intent you also mean the user might ask arbitrary questions or diverge from the script, that is trickier – GraphFlow doesn’t natively support random user interruptions because it expects a defined flow. However, you can mitigate this by including at least one agent (or branch) that catches unexpected inputs. For example, you could have a branch if the user types “help” or something not related to the requested info, which routes to a help agent or simply re-prompts. In practice, since this is a CLI guided dialog, you might not need to handle completely arbitrary queries – the user is mostly following the chatbot’s prompts. If you do need more free-form flexibility at some point, GraphFlow nodes can themselves contain agents that use LLM reasoning to decide how to respond. So you could incorporate a small SelectorGroupChat within a node for a sub-dialog, though that adds complexity. In summary, GraphFlow can cover the expected branches well, and you can design fall-back branches for certain keywords if needed.
- **Human-in-the-Loop Oversight:** GraphFlow integrates naturally with **UserProxyAgent** at points where human input or approval is needed. You can include the `UserProxyAgent` as one of the nodes in the graph (or even have multiple user input nodes). For example, the very first node might effectively be a user node that supplies the username (GraphFlow can start with a user prompt node). Later, if a new user registration requires collecting details, you can have a user node to input those details. Or you can include a confirmation step (e.g. a user node asking “Do you want to proceed with LinkedIn login? [Y/N]”). When GraphFlow reaches a UserProxyAgent node, it will **hand control to the user** and pause until input is given ¹⁷. This is exactly how human oversight is injected: by designing the flow to pause and request feedback at critical junctures. The docs illustrate that in a RoundRobin team, the user agent is called in turn, and in a Selector/GraphFlow, the selection/flow logic determines when the user is called ¹⁸. In

GraphFlow, you explicitly decide where to call the user. This satisfies the requirement for oversight – you have full control over *when* the human gets to intervene.

Keep in mind that when the UserProxyAgent is invoked during a run, it will block the team’s execution until input is provided ¹⁹. The system effectively waits for the human, which is desirable for oversight, but it means you should use these points judiciously. The Autogen guide recommends using such blocking calls only for **short interactions that need immediate feedback** (like a quick approval) ²⁰. In your case, user inputs like entering a username, or confirming a registration, are interactive parts of the normal flow, so it’s fine. Just avoid leaving a user prompt hanging indefinitely; design the UI/CLI such that the user is prompted clearly and expected to respond then and there (which is natural in CLI). Also note that once the flow is paused for human input, you cannot save the team state mid-run easily – but since these are short runs, that’s okay. If a user might *delay* (e.g. not complete registration in one go), a pattern could be to terminate the flow and resume later, rather than keep it blocked (we’ll discuss state persistence later).

- **Modularity & Extensibility:** GraphFlow is highly **modular**. Each node (agent) can be developed and tested in isolation, and new nodes can be added without affecting others except for defining new edges. If you want to extend the chatbot with new modules (like parsing a resume or fetching company info), you can design those as separate sub-flows or additional branches in the graph. For example, after LinkedIn login, you could add a branch: if the next intended task is to parse the user’s resume, have a node for a “ResumeParserAgent” that takes the user’s resume file (which could have been provided earlier or just then). Or you might have a mode where the user can choose to do employer profiling – you could incorporate a decision node (user choice) that leads to an “EmployerProfilerAgent”. Because GraphFlow can do conditional branching, you can even present a menu of next actions to the user and branch accordingly. This is much easier to manage in a graph design than in a rigid round-robin. Each new functionality simply adds some nodes and edges; it doesn’t require rewriting a central loop from scratch. In contrast, with a RoundRobin or basic Selector, adding more agents could complicate turn order or selection prompts significantly. GraphFlow’s explicit structure thus gives you a **scalable architecture** for future growth.
- **Debuggability and Transparency:** A GraphFlow-based implementation will be easier to **debug step-by-step**. Since you dictate the path, you can insert logging at each node transition. Autogen’s `Console(team.run_stream())` can stream the live conversation among agents, which in GraphFlow will correspond to each node’s agent producing a message. You’ll see clearly the sequence: e.g., *UsernameValidatorAgent* -> “Username looks valid”, *DatabaseAgent* -> “User not found”, *RegistrationAgent* -> “Collecting new info”... and so on, in order. There’s no confusion about “why did agent X speak now?” – it’s in the graph logic. Additionally, Autogen v0.4 emphasizes observability: it has hooks for tracing and OpenTelemetry, which you can utilize if needed ²¹. Even without advanced tools, the structured nature of GraphFlow means you can mentally trace the flow like a flowchart during testing. If something goes wrong at a particular node, you know exactly which agent/node is responsible, making it straightforward to isolate issues. This addresses your requirement for tracking agent workflows and debugging incrementally.
- **State and Memory Handling:** By using GraphFlow with defined steps, you implicitly manage state transitions in a controlled way. The conversation state (messages so far) will accumulate in the team as the flow proceeds. If you reuse the team object for multiple runs (or keep the process alive), that context can persist to the next task unless reset ²². This means that if you want to maintain continuity within a session, you can – though for a login flow, you likely reset or terminate once done. As for long-term memory between separate runs (e.g., remembering a user across program executions), Autogen does allow saving the state of agents/teams to disk

²³. In a GraphFlow, you could serialize the entire team (which includes agents and their conversation histories) at the end of a run and reload it later to continue. However, since you are already storing user data in JSON, it might be simpler to just rely on that for persistent user info, rather than trying to persist the chat state. Long-term memory in LLM terms (like remembering detailed conversation context or using retrieval augmentation) can be added later if needed. Autogen supports a `Memory` interface for adding external knowledge into agent contexts ²⁴ – for instance, you could load a user’s profile or past interactions as memory on the assistant agent when a session starts. This is an optional enhancement. The key point is: GraphFlow doesn’t hinder state management; you still have all the underlying capabilities of AgentChat to maintain context or save state. And since state management is not a priority now, you can defer this complexity. Just be aware that if needed, you can **persist agent states** (via `save_state()` on agents/teams) to implement long-term memory or session resume ²³.

Why Not RoundRobin/Selector? To solidify the recommendation, it’s useful to contrast why GraphFlow edges out the simpler options here. RoundRobinGroupChat was tempting for simplicity, but your flow isn’t a strict fixed turn cycle – the number and order of steps vary depending on user input and conditions (for example, if the user is already in the database, you skip the registration step entirely; RoundRobin can’t “skip” an agent without extra hacks). SelectorGroupChat offered dynamic turn selection, which can handle skips, but then you either trust an LLM to pick the right agent (which might be brittle for a structured task), or you essentially code a selection function that does what GraphFlow would do anyway (deciding next agent based on state). In effect, using Selector with a custom function leads you to manually manage state and next-step logic, which is exactly what GraphFlow is designed to streamline – so you might as well use the purpose-built workflow tool. GraphFlow gives you *explicit conditional branching* rather than an implicit one buried in a selection function’s code or an LLM prompt. This clarity is valuable for maintenance and debugging as your system grows. SelectorGroupChat is fantastic for more conversational, open-ended multi-agent scenarios (imagine a brainstorm among agents where it’s not predetermined who should speak when), but for a **guided procedure**, GraphFlow is the more appropriate pattern.

Why Not Magentic-One? As discussed, Magentic-One is powerful but meant for autonomous problem-solving. It introduces an AI-driven planner, which is not necessary when you have a known script. Implementing GraphFlow is much more straightforward for a login flow than configuring an orchestrator agent to “figure out” the login steps on its own. Magentic-One also might execute steps in a less linear fashion (it can interleave tasks, etc.), whereas you want a predictable step-by-step interaction with the user. So while Magentic-One embodies an impressive architecture, it’s not aligned with your immediate needs for simplicity and oversight.

In summary, **GraphFlow with a well-defined directed graph of agents is recommended**. This will let you **modularly implement and test each part of the workflow**, ensure the overall conversation follows the intended path, and still allow user inputs and future extensions. Next, we’ll outline how to implement this incrementally and address debugging, human oversight, and state management in practice.

Incremental Implementation Strategy

Adopting GraphFlow does not mean you have to build everything all at once. You can **incrementally develop and test** the workflow, ensuring each piece works before adding more complexity. Here's a step-by-step implementation plan:

1. **Start with Basic Agents and Tools:** Identify the agents you need and either use Autogen's built-ins or create simple custom agents for each role:
2. *Username Validator / Input Parser:* This could be a simple `AssistantAgent` (with an LLM) that you prompt to check if the provided username meets certain criteria (or is not empty). If you prefer not to use an LLM for this trivial task, you could make a small custom function or a `CodeExecutorAgent` to validate the string. However, using an LLM with a system prompt like "You are a validator agent. Ensure the username is not empty and meets format X..." could suffice, depending on how strict the validation is.
3. *User Lookup Agent:* This might be implemented with a `CodeExecutorAgent` that runs a Python function to search the JSON file for the username. Autogen's `CodeExecutorAgent` can execute code when given in the conversation (if allowed), but an easier approach might be to write a custom agent class that, when prompted with a username, directly checks the JSON and returns a message like "FOUND" or "NOT FOUND". Another approach: you could integrate the JSON lookup as a **tool** and let an `AssistantAgent` call that tool. Autogen v0.4 allows `AssistantAgents` to use tools (the OpenAI function calling or other extension mechanisms) ²⁵. ²⁶ For initial simplicity, a separate code-based agent may be clearer. The key is to ensure this agent can output a signal or content that the GraphFlow can use in a condition (e.g., it might output a message containing a keyword "USER_EXISTS" or "NO_SUCH_USER" which you then use for branching).
4. *Registration Agent:* This would likely be an `AssistantAgent` (backed by GPT-4 or similar) that asks the user for whatever additional info is required for registration (email, etc.), or parses info the user has provided. If registration is complex, you might even break it into sub-steps. But initially, you can make it one agent that formulates the new user entry. It might need to interact with the user (for example, "I need your email to register"). In GraphFlow, you can handle that by having the `RegistrationAgent` output a prompt and then loop to a user input node for the answer, then back to `RegistrationAgent`, etc., or simply have the agent's LLM prompt engineered to ask for all needed info in one go. For incremental dev, perhaps just log a message "Registering new user..." first.
5. *LinkedIn Login Agent:* This is a complex step because it involves Playwright. You have a few options:
 - Use `CodeExecutorAgent` to run a script that opens Playwright and performs login (if the environment allows that). You might have to ensure the agent has access to necessary libraries (Autogen's extension for web, or a custom integration).
 - Write a **custom agent** (subclass `BaseAgent` or `AssistantAgent`) that overrides its `step()` to perform the Playwright login action behind the scenes when triggered, then returns a success/failure message. This might be cleaner, as you can encapsulate the automation logic in Python code rather than LLM. The agent could even take the username (and maybe password, though password should be handled securely) from context and do the login.
 - Alternatively, treat Playwright as a "tool" and have an LLM agent invoke it. Autogen does have a web-surfer extension (for browsing web), but logging in to LinkedIn might not be directly covered. A custom approach is likely needed. For the first iteration, you can stub this out (simulate a login success message without actually doing it) just to get the flow working. Later, integrate the real Playwright code.

6. *Credential Storage Agent*: Possibly not an agent at all – once login is successful, you might just call a function to encrypt and save credentials. This can be done inside a CodeExecutor agent or even inside the LinkedIn agent's custom code (i.e., after a successful login, it stores the credentials). To keep design modular, you could make a separate small agent responsible for saving data (or just use a final step in GraphFlow that calls a function).
7. *UserProxyAgent*: This is needed for any point you want actual user CLI input. At minimum, you'll use it at the start to get the username. You might also use it during registration (to get new user info) or for confirmations. Create a `UserProxyAgent("user", input_func=input)` so it reads from console ²⁷. For testing, you might simulate user input or just be ready to type responses.

As you implement each agent, **unit test them individually** if possible. For example, you can call `await agent.run(task="some input")` on an AssistantAgent or custom agent to see what it outputs given a prompt, or ensure the CodeExecutor can access the JSON, etc. This isolates issues early.

1. **Define a Simple Flow Graph**: Using Autogen's `GraphFlow` API, start by constructing a very basic directed graph connecting a few of these agents in order. The Autogen docs suggest using `DiGraphBuilder` to add nodes and edges fluidly ²⁸. For instance:
 2. Node1: UserProxyAgent (get username) -> Node2: ValidatorAgent (validate format).
 3. For now, connect Node2 unconditionally to Node3: LookupAgent (check JSON).
 4. Then connect Node3 to Node4: perhaps a DummyAgent that just prints "(End of demo)". This is just to test the linear flow.

Leave branching for later; first confirm that you can run the GraphFlow from Node1 through Node4 in sequence. Use `Console(team.run_stream())` to run it and manually input a username at the prompt. You should see each agent's turn in order. This verifies the basic GraphFlow wiring and that agents can produce and consume messages.

1. **Add Branching Logic**: Next, enhance the graph with a branch for user existence:
2. For Node3 (LookupAgent), set up two outgoing edges:
 - If the user exists (the agent might output a message containing "FOUND" or some flag in content/metadata), go to Node5: LoginAgent.
 - If the user is not found, go to Node6: RegistrationAgent.
3. Then decide how the flow should converge or proceed:
 - After RegistrationAgent completes creating the new user, you might then proceed to the LinkedIn LoginAgent (Node5). So Registration node would connect to Login node (Node5) as well, but only after it finishes the registration process.
 - After LoginAgent (Node5) succeeds, go to Node7: CredentialStorage/Session agent (or simply terminate if that agent handles final logging).
 - If LoginAgent fails (perhaps multiple attempts or immediate failure), you could loop back or go to an error handler. Initially, you might skip this for simplicity or just terminate on failure with a message.
4. To implement conditions on edges in GraphFlow, you'll likely use a lambda or predicate that inspects the last message from the source agent. Autogen might allow matching text (e.g., `if "FOUND" in message.content`). Alternatively, you can design the LookupAgent to return a structured message like `{"user_exists": true}` in metadata that your edge condition reads. Simpler: have it output exactly "USER_FOUND" or "USER_NOT_FOUND" as the content, and key off that.
5. Similarly, for the LoginAgent, you can have it output "SUCCESS" or "FAIL" and branch accordingly (e.g., success leads to storing session, fail could loop to ask for credentials again or to a human intervention).

This step introduces the main **conditional logic**. Test it by trying a username that you set up in the JSON (to follow the exists->login path) and one that is not in JSON (to follow the new user path). At first, you could have the RegistrationAgent simply output a dummy “User registered” message without actually gathering details, just to drive the flow forward to login. We will integrate actual user inputs in the next step.

1. **Integrate Additional User Inputs (Human-in-the-Loop):** Now incorporate the `UserProxyAgent` at any points where you need to gather input or approval mid-flow:
2. Already at Node1 we have it for the username. You might also need it after the RegistrationAgent prompts for details. One approach is to break the registration interaction into multiple nodes:
 - Node6: RegistrationAgent asks “Please provide your email and name.”
 - Node7: UserProxyAgent (user provides those details).
 - Then loop back to Node6: RegistrationAgent (which now sees the user’s response in context, and can create the account or ask for more info if needed).
 - This loop continues until RegistrationAgent is satisfied (maybe just one cycle if you gather everything at once).
3. Alternatively, you can handle multi-turn interaction within a single agent by giving it the entire conversation (the agent would output a message asking for info, then on next turn the user replies, etc.). However, in GraphFlow you typically want each node to represent a single agent action. So modeling it explicitly (agent -> user -> agent) is cleaner.
4. If you want a confirmation step before proceeding to login (for instance, “User registered successfully. Press Enter to proceed to LinkedIn login.”), you can insert another user node there.
5. Consider an oversight scenario: imagine after registration and before performing the automated login, you want a human to confirm (maybe the user needs to confirm they want to use their LinkedIn credentials). You can put a `UserProxyAgent` node that simply asks “Type YES to continue to LinkedIn login.” If the user types yes, proceed; if no or timeout, you could terminate. This is an example of human-in-loop control that GraphFlow can enforce by halting until input is given.
6. With each added user interaction, test the flow. Keep in mind from the docs that in Selector/GraphFlow, the selector or flow logic decides when to call the user agent ¹⁸ – here that logic is your graph design. You must ensure the user agent node has an incoming edge from the correct previous node. Once triggered, it will wait for input and then output the user’s message as a `ChatMessage` to the next node.
7. **Implement Agent Behaviors and Tools:** At this stage, the structure is in place. Now focus on the **meat of each agent’s behavior**:
 8. Fill in the actual JSON lookup code in the lookup agent (or integrate it properly as a tool call). Ensure it can read from the JSON file path (consider security, but since this is local CLI, it’s fine).
 9. Improve the username validation agent’s prompt or code so it catches empty input or disallowed characters and appropriately asks again (e.g., output a specific flag if invalid to trigger the loop).
 10. Develop the registration agent to actually store new user info in the JSON (or at least prepare the data for it). Possibly the agent itself can call a function to append to JSON once it has the details. If using an LLM, beware of trusting it to format JSON correctly; you might again use a code tool or a template.
 11. Integrate the LinkedIn Playwright automation. This will likely be the most challenging part technically (outside Autogen’s scope). You might create a function like `attempt_linkedin_login(username, password) -> bool` and then call it from a `CodeExecutorAgent` or a custom agent’s code. Since passwords are involved, ensure you handle

them securely (you might prompt the user for a password through the UserProxyAgent rather than store it in plain text anywhere). Possibly, the user enters their LinkedIn password at runtime and you pass it to the LinkedIn agent node which then logs in. In any case, test this outside the flow first (make sure Playwright can run from CLI and do a headless login).

12. For credential storage, implement encryption (maybe using a library or a simple symmetric encryption for now) and save to a file or database. This could also be a function call triggered after login success. You might incorporate it into the LoginAgent (post-login) or have a separate agent node solely for “StoreCredentialsAgent” that takes the login result and performs storage.
13. **Termination conditions:** Even though GraphFlow will naturally end when it runs out of nodes to execute (i.e., when you design the graph's end), it's good to also set a termination condition. For instance, you can use `TextMentionTermination("exit", sources=["user"])` to allow the user to type "exit" to abort ²⁹. Or use `MaxMessageTermination` or `ExternalTermination` if you want a manual cancel. Given it's CLI, a simple “type exit to quit” (monitored via the user agent input) is helpful. Termination conditions can be combined (ORed) to include multiple triggers ³⁰. Ensure to reset termination if you reuse the team in a loop.
14. **Testing Workflows Step-by-Step:** As you implement each piece, perform tests:
 15. Test the **existing user path**: Does the flow correctly bypass registration and go straight to login? Is the conversation coherent (e.g., it shouldn't ask for details if user exists)? Does the login step get the needed credentials (you might have stored the password earlier or ask for it on the spot)?
 16. Test the **new user path**: Try a username not in JSON. The validator -> lookup should detect not found, then go to registration. Step through registration: input details as prompted. After providing details, ensure a new entry is created (or at least logically, that the flow marks the user as now “exists”). Then it should proceed to login. If you haven't built storing to JSON yet, at least confirm the flow continues to login with whatever data it has.
 17. Test **invalid inputs**: e.g., user presses Enter with no username, or enters an invalid format. The ValidatorAgent should catch it and ideally loop back to ask again. Make sure your graph's loop on invalid username works (you might simulate multiple loops).
 18. Test **user oversight steps**: Try saying “no” or something at a confirmation step (if you added one) to see if the flow terminates or behaves correctly. Also test the termination keyword (“exit”) to confirm the run stops.
 19. Use the console logging or add printouts to confirm each node triggers in the right order. If something goes wrong, the GraphFlow might throw if it hits a node with no outgoing edge and not terminated, etc., so watch for that.
20. **Debugging tip:** Autogen's `team.run()` will return a TaskResult with the full conversation. You can examine it to see all messages and ensure the content tags for conditions are present. For live observation, `run_stream` with `Console` is very convenient, as it prints each message with the agent's name ³¹.
21. **Iterate and Refine:** With basic functionality working, refine the system:
 22. Improve prompts for better user experience (make the assistant's messages clear and user-friendly).
 23. Add error-handling nodes. For example, if LinkedIn login fails, maybe allow a second attempt (loop back to ask for password again) and after 2 failures, either give up or ask if they want to skip. GraphFlow can implement a **loop with a counter** by keeping count in an agent's state or

using a Loop with an exit condition in the DiGraphBuilder (the docs mention loops with safe exits ³²).

24. Integrate memory or context if needed: e.g., if in the future you parse a resume, you might store it in a memory vector and retrieve info later. For now, ensure that any important data (like the newly created user record or session token) is passed along. You might attach it to the conversation context (e.g., the RegistrationAgent could output the new user's ID which the LoginAgent then uses). Alternatively, share data via an external variable or the JSON itself (after adding the user, the login agent could just look them up from JSON for credentials).
25. Ensure **security** for credentials: You mentioned encryption – implement that before finalizing. Also be careful not to log sensitive info in plain text during debugging.
26. Perform an end-to-end test where you simulate a real scenario: an unknown user signs up and logs in, then run again with the now-registered user to see the existing-user path.

Throughout this incremental build, GraphFlow's explicitness helps in adding one piece at a time. For instance, you can first connect registration to login directly (with perhaps a dummy password) before worrying about asking the actual password from the user. Once stable, you add a user input node for the password prompt. Each addition is local to a part of the graph.

Finally, since GraphFlow is new in v0.4, **keep an eye on the Autogen documentation and examples**. The official examples (like travel planning, company research, etc.) might demonstrate GraphFlow usage or patterns of agent interactions that you can adapt ³³ ³⁴. Given it's experimental, test carefully and consult the docs if something doesn't behave as expected. The good news is that your use case is exactly the kind of **deterministic multi-step process** GraphFlow was built to handle, so it should serve you well.

Debugging and Testing Workflows

Ensuring the system works correctly is as important as building it. Autogen v0.4 provides tools for **observability, logging, and interactive debugging** that you should leverage:

- **Console Streaming:** We've mentioned using `await Console(team.run_stream(...))` during development. This will print each message exchange to your console with the agent names labeled ³¹. It's extremely useful for watching the conversation unfold in real time. You will see exactly when the `UserProxyAgent` pauses for input, what each agent outputs, and when termination triggers. This immediate feedback makes it easier to pinpoint logic errors (e.g., if a branch condition didn't fire as expected, you'll see the wrong agent speaking next).
- **Logging and Tracing:** Autogen supports integration with logging frameworks and OpenTelemetry for tracing agent events ²¹. For a CLI app, you might not need full telemetry, but you can certainly add simple logging (using Python's `logging` module or prints) inside custom agent code or in the selection functions/edge conditions to record decisions. For instance, if you wrote a custom selector function (in case you tried SelectorGroupChat) you could log which agent was chosen. In GraphFlow, you can log when a conditional edge is taken ("LookupAgent decided user exists, going to LoginAgent"). This kind of logging, even if to stdout, can greatly help during debugging complex flows or user issues reported later.
- **Step-by-Step Inspection:** If something isn't working, break down the problem. Because your design is modular, you can test a subset of the flow by modifying the start or end of the graph. For example, if the registration part is failing, you can test the RegistrationAgent and subsequent nodes in isolation by crafting a small graph that starts at registration (feeding it a context of a given username). Autogen might allow starting a GraphFlow at an intermediate node or you can

simulate it by calling the agent's run method manually with a dummy prompt. Use the **AgentChat API in interactive mode** (like a notebook or REPL) to call agents one by one if needed.

- **Human-in-the-loop Testing:** To test human oversight mechanisms, simulate a non-cooperative or confused user to see how the system responds. For instance, enter an obviously invalid username and see if it correctly reprompts. Or at a yes/no confirmation, type something unexpected like “maybe” – does the system handle that gracefully (perhaps by treating anything not “yes” as a “no”)? You might need to refine prompts or add slight logic to handle such cases.
- **Termination Conditions:** Explicitly test the termination paths. If you set a keyword termination (like “exit”), try typing it at various points (e.g. at the username prompt, or during registration) to ensure the conversation stops. Also test the normal end of the flow – the run should terminate when the final step is done. Verify that the `TaskResult` is returned and that resources (like the Playwright browser or model clients) are closed properly in a `finally` block (as shown in Autogen examples) ³¹.
- **Autogen Studio (Optional):** Microsoft provides **AutoGen Studio**, a low-code GUI for prototyping agent workflows ³⁵. In your case, since you're coding directly, you might not need it. But know that it exists – it can visualize agent interactions and might be an interesting way to **visualize the GraphFlow** if supported. It's mentioned that Studio has real-time agent updates and message flow visualization ³⁶. This could be useful if you want to see a diagram of your GraphFlow execution in real-time, which might help in debugging or demonstrating the flow to others. It's not necessary for development, but a good example of how Autogen's ecosystem supports debugging.

In short, adopt a **test-as-you-go approach**. Each agent and each branch added should be immediately tested in isolation or in small groups. By the time you have the full end-to-end flow, you'll have high confidence in each component. Leverage Autogen's built-in messaging logs and keep your eye on the console outputs to verify the sequence of events. Because GraphFlow is deterministic, a bug will usually be reproducible, making it easier to trace and fix.

Human-in-the-Loop Considerations

Having a human in the loop (in this case, the end-user providing inputs or confirmations via the CLI) is a core part of your chatbot. Here are some additional considerations to effectively integrate human oversight while using the Autogen framework:

- **UserProxyAgent Placement:** We've already covered where to put `UserProxyAgent` nodes in the GraphFlow. Generally, any time you need the user's input or approval, that's a spot for a `UserProxyAgent`. Design the **prompts carefully** – since this is CLI, the user might not have a GUI button, so they must type responses. Make it clear what they should input (e.g., “Enter your LinkedIn password (or type 'skip' to cancel): ”). The `UserProxyAgent` will display this prompt if you include it in the preceding agent's message.
- **Blocking Nature and User Experience:** Remember that when waiting for user input, your program is essentially paused (the event loop waits). That's fine for interactive mode, but if you ever consider running this in a web context, you'd handle it differently (like using `HandoffTermination` or external signals). In CLI, just ensure the user knows to respond. If there's a possibility the user might not respond (walks away, etc.), consider implementing a timeout.

Autogen might not have a built-in timeout on `UserProxyAgent` input by default (it just waits indefinitely on `input()` by default). You could customize `input_func` to abort after some time or number of retries. Alternatively, provide a way to exit (like typing “exit”) as mentioned. The **HandoffTermination** mechanism ³⁷ is another pattern: an agent can output a special handoff message that triggers termination so the external application (your script) can take over. For example, an agent could detect no response and issue a handoff to end the run gracefully, rather than hanging. This is advanced usage; for now, a simpler approach is fine given a direct CLI context.

- **Human Oversight vs. Automation:** Decide how much autonomy you want to give the agents before requiring human confirmation. For instance, should the bot automatically register a new user, or should it ask “Shall I create an account for username X with email Y?” and wait for the user to confirm? In a high-stakes domain, you’d ask; here, it’s probably fine to just do it if the user provided the info. But for something like attempting a LinkedIn login, maybe you want a quick “Ready to attempt LinkedIn login? (y/n)” confirmation, especially since it might trigger external browser automation. Those design choices are up to you, but GraphFlow allows either approach. Keep in mind the Autogen guidance that blocking interactions should be for **immediate feedback** ²⁰ – asking for confirmation right before doing a potentially risky action fits that description.
- **User Interruptions:** Even though the flow is designed, users might still deviate (entering unexpected input). For example, the user might type “help” instead of a username. If you want to handle that, you could pre-empt it by documenting usage (“Type ‘exit’ to quit, or follow the prompts.”). If needed, you could include a special condition on the username input node: if input == “help”, route to a `HelpAgent` that prints some help text, then loop back to ask username again. This kind of out-of-band request can be handled by additional branches.
- **Transparency and Control:** One aspect of human-in-the-loop is making the system’s actions transparent to the user. Since you have multiple agents working behind the scenes (some possibly using LLMs or tools), it’s good to log or announce certain steps to the user. For instance, when transitioning to the LinkedIn login step, the bot might say “Okay, I will now attempt to log in to LinkedIn with the provided credentials.” This keeps the user informed and feels like oversight even if they aren’t actively intervening at that moment. It’s generally a good practice in agentic systems to be clear about what the AI is doing or about to do, especially when using tools like web browsers.
- **Iterative Improvement with Human Feedback:** As a final note, the human user can also be a source of feedback to improve the flow. If you find during user tests that people are getting stuck or confused at a certain step, you can adjust the flow or add an agent to handle that. The modular structure makes such adjustments relatively isolated (e.g., if users often input an unrecognized command at some point, you can add a branch to catch it without overhauling the whole system).

In essence, the human-in-the-loop design in Autogen (via `UserProxyAgent` and orchestrated via GraphFlow) gives you a **high degree of control** over when and how the user participates in the conversation. By thoughtfully placing those interactions and handling their responses, you ensure that the user is neither overwhelmed with too many inputs nor sidelined when their confirmation is needed. This fulfills the oversight requirement by making the user an integral part of the loop at key moments, rather than a passive observer of an autonomous process.

State Management and Long-Term Memory (Optional)

While not a priority for your immediate goals, it's worth discussing how you might handle state and memory if the need arises – especially as you consider long-term usage or session continuity.

- **Conversation State Persistence:** By default, an Autogen team (whether RoundRobin, Selector, or GraphFlow) will retain the conversation history in memory as long as the team object exists. That means if you call `team.run(task=". . .")` multiple times sequentially, the second run will still have context from the first (unless you reset). The SelectorGroupChat documentation notes that *"Once the team finishes the task, the conversation context is kept within the team and all participants, so the next task can continue from the previous context"* ²². This likely applies to GraphFlow as well, since it's a kind of team. For a CLI chatbot scenario, typically you might run one GraphFlow per user session. If you wanted to allow a user to do multiple actions in one session (like after login, then do another task without restarting), you could indeed continue using the same team to maintain memory of what happened. In your current design (which mostly ends after login), you might not need that. But if you add more post-login features (like "now that you're logged in, do you want to parse your resume?"), keeping the context around (so the agent knows the user's name, etc.) can be convenient. You can always call `team.reset()` to clear history if you want to start fresh for a new user or a new session.
- **Saving/Loading State Between Runs:** If you want long-term memory (e.g., the next time the user runs the program, it remembers something from before), you have to persist state to disk or database. Autogen v0.4 provides a mechanism to **save the state of agents, teams, and termination conditions** to a Python dict or file ³⁸ ³⁹. For example, each `AssistantAgent` has a `save_state()` method that captures its internal state (conversation messages, memory, etc.) ⁴⁰. Similarly, the Team (GraphFlow) might have a `save_state`. You could serialize these (maybe via `pickle` or JSON if supported) at the end of a session and reload them later to restore the conversation. This is particularly useful in web apps where the backend might be stateless between requests ⁴¹. In a CLI app, you could simply keep the program running to maintain state, but if that's not possible, you can persist needed info.

In practice, you might not want to reload a whole conversation history every time – instead, you might just extract key pieces to store. For instance, after registration, you store the new user's credentials in your JSON database (which you already plan to do). That's effectively long-term memory outside the chat. If there were other details (like the conversation or decisions made), you could log them or store them if needed. Another example: if using an LLM agent, it might accumulate a lot of tokens in the prompt after many turns. Saving state and reloading could allow you to start a new run with a summary of the old conversation rather than the full log (though Autogen doesn't auto-summarize, you'd have to do it).

- **Memory Extensions (Knowledge Bases):** Autogen has a concept of memory stores for retrieval augmented generation (RAG) ²⁴. While this might be beyond the scope of your login flow, it could come into play for future modules. For example, if you have a resume parsing module, you might store extracted facts from the resume in a memory (vector database) so that the assistant agent can recall them when answering questions or filling forms. The Memory API allows you to `add` data and later `query` relevant data to insert into an agent's context ²⁴. This could serve as a long-term memory of user-specific info that persists across sessions (since you could save the memory DB on disk). Again, for a straightforward login system, this is probably overkill – your JSON user database and some logs might suffice as "memory". But it's good to know the framework supports this pattern.

- **Session Tokens or Cookies:** Since you mentioned saving login sessions for future use, consider how you'll store session cookies or tokens retrieved after LinkedIn login. If your LinkedIn login agent can output the session token, you should store that securely (perhaps in your JSON store or another secure file). That's not exactly Autogen "state" but rather application state. You might encrypt it as planned. The next time the user logs in via the chatbot, you could detect they have a saved session and skip the Playwright login entirely (which would be a nice optimization/feature). That would introduce another branch: existing session -> use it, else -> do login. This again can be handled in GraphFlow by adding a check at the appropriate point (maybe part of user lookup, or separate check after login attempt fails).
- **Cleanup and Resource Management:** When dealing with long-term state, also plan for cleanup. If your GraphFlow or agents maintain open resources (like an open browser or file handles), ensure they are closed at termination. For example, Autogen's examples show closing the model client and web surfer agent after the team run ⁴². Similarly, you should call `await web_surfer.close()` or any such method if using those tools. For Playwright, you'd close the browser context. If you keep the program running waiting for user input at multiple stages, make sure there are no stray resources left open between stages. Proper use of `finally` blocks or context managers helps here.
- **Edge Cases for State:** If a user stops mid-process (say, they close the CLI after entering username but before completing login), you might end up with half-registered data. Think about how you'd handle that on next run – perhaps detect an incomplete registration and offer to resume. GraphFlow could allow resuming if you saved the state, but it might be simpler to just start over or have a checkpoint system in your JSON (like a flag "email collected, pending LinkedIn login"). These design decisions ensure your system is robust in real-world use.

In conclusion on state: **You can achieve long-term memory if needed, but it requires explicit handling.** Autogen's ability to save agent/team state ²³ is handy for carrying over conversation context, but in a scenario like yours, it might be sufficient to persist the essential data in your own database and simply re-initialize a new conversation next time using that data. If later you build more conversational features where the history is important to keep (like remembering previous queries or preferences), you can use the provided state management functions or memory stores to maintain continuity. Since it's not a pressing requirement now, the best approach is to design your system such that adding persistence later will not be too difficult (e.g., keep user-specific info in easily accessible structures). The modular nature of GraphFlow and Autogen agents will help – because each agent does one thing, you know exactly what state to capture for each (for instance, the RegistrationAgent might produce a `user_profile` object – you can decide to save that externally after it runs, etc.).

Conclusion and Next Steps

By choosing **GraphFlow as the backbone architecture**, you gain a structured, transparent way to orchestrate your multi-agent chatbot. This addresses the need for a clear, mostly-linear flow with controlled branches, making the system **easier to debug and extend**. The use of **UserProxyAgent** at key steps ensures the user remains in control where necessary, fulfilling the human-in-the-loop requirement. Other patterns like RoundRobin or Selector, while simpler, would become unwieldy as the logic branches out. The Magentic-One orchestration approach, while powerful, is excessive for a known sequence of actions and would sacrifice some clarity and oversight.

With the recommended approach, you can confidently implement the current login and registration features, and later integrate additional capabilities (resume parsing, etc.) by simply **growing the graph**

of interactions. Each new agent or tool can be plugged in as a node with minimal impact on the rest of the system, preserving modularity. Debugging is facilitated by the step-wise nature of GraphFlow, and testing each part in isolation is straightforward. Should the need arise, Autogen's features for memory and state can be layered on to provide continuity across sessions, and the robust termination and messaging system can handle user interrupts or errors gracefully.

In moving forward, make sure to leverage the Autogen documentation and examples for GraphFlow and related features – they provide guidance and patterns that can save you time (for example, how to structure conditional logic, or how to integrate custom tools). Keep your implementation **iterative**, as outlined, to continually validate each piece of the puzzle.

By following this structured approach, you will build a chatbot that not only meets the current requirements but is **architecturally prepared** for future enhancements. Good luck with your implementation, and enjoy the process of assembling your agent team with Autogen v0.4's powerful toolkit!

Sources:

- Microsoft Autogen v0.4 Documentation – *Teams and GroupChat Patterns* 1 2 43 15
- Microsoft Autogen v0.4 Documentation – *GraphFlow (Workflow) Guide* 13 11
- Microsoft Autogen v0.4 Documentation – *Magentic-One Overview* 6 8
- Microsoft Autogen v0.4 Documentation – *Human-in-the-Loop and UserProxyAgent* 44 17 20
- Microsoft Autogen v0.4 Documentation – *State Management and Memory* 22 23 24
- Microsoft Autogen Blog – “Building Agentic Solutions with Autogen 0.4” (Design patterns and stages) 45 46
- Microsoft Autogen GitHub – *README and Quickstart Examples* (multi-agent usage and termination) 47 48

1 30 Teams — AutoGen

<https://microsoft.github.io/autogen/stable/user-guide/agentchat-user-guide/tutorial/teams.html>

2 3 4 5 22 Selector Group Chat — AutoGen

<https://microsoft.github.io/autogen/stable/user-guide/agentchat-user-guide/selector-group-chat.html>

6 7 8 Magentic-One — AutoGen

<https://microsoft.github.io/autogen/stable/user-guide/agentchat-user-guide/magentic-one.html>

9 10 11 12 13 14 28 32 33 34 GraphFlow (Workflows) — AutoGen

<https://microsoft.github.io/autogen/stable//user-guide/agentchat-user-guide/graph-flow.html>

15 16 24 25 26 37 43 45 46 48 Building Agentic Solutions with Autogen 0.4 | Microsoft Community Hub

<https://techcommunity.microsoft.com/discussions/azure-ai-services/building-agentic-solutions-with-autogen-0-4/4394757>

17 18 19 20 27 44 Human-in-the-Loop — AutoGen

<https://microsoft.github.io/autogen/stable/user-guide/agentchat-user-guide/tutorial/human-in-the-loop.html>

21 35 36 AutoGen reimagined: Launching AutoGen 0.4 | AutoGen Blog

<https://devblogs.microsoft.com/autogen/autogen-reimagined-launching-autogen-0-4/>

23 38 39 40 41 Managing State — AutoGen

<https://microsoft.github.io/autogen/stable/user-guide/agentchat-user-guide/tutorial/state.html>

29 31 42 47 GitHub - microsoft/autogen: A programming framework for agentic AI PyPi: autogen-agentchat Discord: <https://aka.ms/autogen-discord> Office Hour: <https://aka.ms/autogen-officehour>
<https://github.com/microsoft/autogen>